

Bridging the Gap in Text-Based Emotion Detection

Sarvnazsadat Roumianfar, Ali Abbasi
Politecnico di Torino, Turin, Italy

CONTENTS

I	Introduction	1
II	Problem Definition	1
III	Proposed Solutions	1
III-A	Transformer-Based Model Fine-Tuning (encoder-only)	1
III-B	Loss Function Adjustment with Class Weights	1
III-C	Transformer-Based Model Fine-Tuning (decoder-only)	1
III-D	Layer freezing	2
III-E	Lightweight Fine-Tuning with LoRA and Qunatization	2
IV	Criteria for Assessing Solutions	2
V	Research Methodology	2
VI	Analysis and Interpretation	3
VII	Conclusions and Recommendations	3
	Appendix A: Model Training Configurations and Performance	4
	References	4

LIST OF FIGURES

1	Emotion distribution in the English subset of the dataset, illustrating class imbalance. Some emotions, such as fear, are significantly overrepresented, while others, such as anger are underrepresented (disgust class is absent for english dataset).	1
2	Training and evaluation metrics for the DeBERTa model. While training loss steadily decreases, evaluation loss begins to increase after epoch 5, indicating overfitting. Similarly, evaluation F1 scores (both micro and macro) fluctuate rather than improve, suggesting that DeBERTa did not yield notably satisfying results on our task.	2
3	LLaMA demonstrates a better leaning curve than DeBERTa and BERT, indicating that it has a better learning process and better fits to our problem	3
4	Comparison of macro and micro F1 scores for BERT, DeBERTa, and LLaMA with Unfreezed Layers models.	4

LIST OF TABLES

I	Performance Comparison of LLaMA Variants on Emotion Classification	3
II	Performance Comparison of unfreezing different layers of LLaMA model	4
III	Training Configurations and Evaluation Results	4

Bridging the Gap in Text-Based Emotion Detection

Abstract—This study presents a fine-tuned large language model for multi-label emotion classification based on the BRIGHTER dataset, as a part of the SemEval 2025 Task 11, Track A. The purpose was to develop a transformer-based classifier capable of detecting multiple emotions simultaneously from short social media posts. Due to the shortage of resources that we had, we tried to adapt to solutions that need less GPU and memory resources while keeping our performance comparable to other result in the competition. Our final model achieved a macro F1 score of 0.7114 on the English subset of the BRIGHTER dataset, outperforming all leaderboard submissions to date.

I. INTRODUCTION

Emotion classification in text is challenge due to the complexity of human expression. Accurately assigning a label to a sentence or phrase is difficult because individuals often express their emotions in abstract way or use implicit meanings. As part of the SemEval-2025 Task 11 competition [6], this project aims to develop large language model (LLM)-based architectures capable of identifying multiple emotions within English-language texts. We use the BRIGHTER dataset [5], a multilingual benchmark for this task, which includes emotion-annotated texts in 28 languages. This report outlines the evolution of the project from traditional pretrained transformers such as BERT*** and DeBERTa*** to more advanced LLMs like LLaMA***, which ultimately yielded the highest performance.

II. PROBLEM DEFINITION

The problem is to design a model capable of detecting multiple concurrent emotional states from short text inputs. SemEval-2025 Task 11 (Track A) defines this problem in a monolingual setting, where the goal is to predict multiple binary labels across 6 emotion categories: joy, sadness, anger, fear, surprise, and disgust (While the disgust label is not present for the English subset of dataset and they are all null). Each label is independently annotated, and the model must learn to correctly identify all applicable emotions for a given sentence. The challenge includes:

- Semantic ambiguity, where phrases can evoke different emotions based on the context.
- Class imbalance, where some emotions (e.g., fear) are overrepresented and others (e.g., anger) are rare (see Figure 1).
- Training efficiently under compute constraints, where Full fine-tuning of large language models is resource-intensive. Therefore, we had to apply some techniques like quantization [4] to reduce model size and memory usage and parameter fine-tuning methods like LoRA [3] to train effectively with limited hardware resources.

Given these challenges, the primary objective of this project is to develop a robust and scalable LLM-based solution that

accurately predicts multilabel emotional states from English text.




Fig. 1. Emotion distribution in the English subset of the dataset, illustrating class imbalance. Some emotions, such as fear, are significantly overrepresented, while others, such as anger are underrepresented (disgust class is absent for english dataset).

III. PROPOSED SOLUTIONS

A. Transformer-Based Model Fine-Tuning (encoder-only)

We explored several transformer-based models. Since, multilabel emotion Detection is a Classification problem, Encoder-Only Models are the best choice for this kind of Tasks. Our early attempts involved fully fine-tuning pretrained light models like BERT and DeBERTa

B. Loss Function Adjustment with Class Weights

To handle the effect of imbalanced labels, we computed class weights for each emotion label in the training data, so we managed to improve the model's sensitivity to underrepresented emotions.

C. Transformer-Based Model Fine-Tuning (decoder-only)

Since we didn't get satisfying result from BERT and DeBERTa (Figure 2) we were looking for other models that are more powerful and larger scale. LLaMA was our choice but they were 2 problems with LLaMA: First, LLaMA being a decoder-only model, is not naturally suited for classification task. However, It can be adapted for multi-label classification by adding a classification head. Second problem is that It is impossible to fully fine-tune llama due to our limited resources for this large scale model. The solution for that is freezing most of the layers of that.

D. Layer freezing

To balance performance with computational efficiency we used layer freezing. Layer freezing means preventing some layers of a model from updating during training and just fine-tune the last layers. We often freeze the first layers because they capture general, low-level features (like word or token patterns), and unfreeze the last layers to let the model learn task-specific representations without forgetting the core language knowledge. We have experimented with freezing the early layers of the model and unfreezing different layers of LLaMA. The results showed that unfreezing last 4 layers has brought us the most effective and most accurate one.

E. Lightweight Fine-Tuning with LoRA and Qunatization

In order to train under limited computational resources, we also applied Low-Rank Adaptation (LoRA) to introduce task-specific trainable parameters while keeping the majority of the model frozen. Additionally, we employed **4-bit quantization** to compress the model weights and reduce memory consumption in training.

Finally we do the hyperparameter tuning by just changing hyperparameters like batch size, number of epochs, learning rate, etc., to find the best-performing model




Fig. 2. Training and evaluation metrics for the DeBERTa model. While training loss steadily decreases, evaluation loss begins to increase after epoch 5, indicating overfitting. Similarly, evaluation F1 scores (both micro and macro) fluctuate rather than improve, suggesting that DeBERTa did not yield notably satisfying results on our task.

IV. CRITERIA FOR ASSESSING SOLUTIONS

To evaluate the performance of our models, we consider 3 metrics below:

- Micro F1-Score looks at all predictions together. It counts all correct and incorrect predictions across all classes, showing how well the model performs overall.
- Macro F1-Score treats each class equally. It calculates the F1-Score for each class separately and then averages them. This helps us to see how good the model performs on common and rare emotions.
- Training and validation loss per epoch to observe the behavior of the model in each epoch to detect Overfitting (E.g. Fig. 2), Underfitting, and the right balance between them. They can also help us to decide when to stop training or still the model can be trained without overfitting.

V. RESEARCH METHODOLOGY

We focused on building a multi-label emotion classification system using the English subset of the BRIGHTER dataset. Our methodology was based on three main considerations: (i) relevance to real-world emotion recognition tasks, (ii) alignment with our evaluation criteria (macro and micro F1-score), and (iii) practical feasibility given our computational resources.

- **Dataset:** Our experiments are based on the BRIGHTER-emotion-categories dataset, published as part of SemEval-2025 Task 11. The dataset consists of short social media posts annotated with multi-label emotion tags. Each instance can include one or more of five core emotions: joy, sadness, fear, anger, and surprise. Labels are represented using a multi-hot encoding scheme to represent the multi-label structure of the task. BRIGHTER includes data in 28 languages; however, For this submission, we used only the English subset ("eng"), which includes:
 - Training set: Multilabel annotations in binary format.
 - Development set: A validation split is used for model tuning and early stopping, with gold labels.
 - Test set: Used for final evaluation, with gold labels that are provided for offline scoring.
- **Data Preprocessing:** Emotion labels were converted as binary vectors of length five to support multi-label classification. Then the dataset was tokenized using a pretrained tokenizer, and split into training, validation, and test sets.
- **Model Selection and Baselines:** All the models that we have selected, are based on the transformer architecture, that uses multi-head self-attention mechanisms, feed-forward neural networks, residual connections, and layer normalization. We chose BERT, DeBERTa, and LLaMA that are all pre-trained models, because of their great performance in different NLP tasks. They allowed us to compare encoder-based and decoder-only structures. Each model was adapted for multi-label classification by attaching a classification head with five output neurons and applying a sigmoid activation function. Binary cross-entropy was used as the loss function.
 - 1) BERT (Bidirectional Encoder Representations from Transformers): is based on the transformer encoder and captures context from both directions of a sentence. We used the BERT-base version with 12 layers and performed full fine-tuning on all parameters. The inputs were tokenized with the original BERT tokenizer and fed through the model. We have trained with batch size of 16 in epochs of 5, 15 and 20 in different runs to evaluate the performance of stability. [1].
 - 2) DeBERTa (Decoding-enhanced BERT with Disentangled Attention) by Microsoft: is based on the transformer encoder like BERT, but it improves

performance by separating word content and position in the attention mechanism. It includes an enhanced decoder to predict masked words better. We used DeBERTa-base (12 layers) and applied full fine-tuning across all layers. We have trained the model with a batch size of 16 across 5, 10, and 15 epochs. In one of our experiments we have applied the learning rate of 2×10^{-5} and a weight decay of 0.01. [2]

- 3) **LLaMA (Large Language Model Meta AI):** is a decoder-only transformer architecture developed by Meta. We used the smaller version, Llama-3.2-1B, which contains 1.2 billion parameters and 16 transformer layers. We have tested batch sizes of 2, 3, 6, and 8 across 5, 8, 10, and 15 epochs. Smaller batch sizes were used to fit within the memory limitations, and larger batch sizes was because of improving training stability. We used a learning rate of 1×10^{-5} and a weight decay of 0.01. Lower learning rate was used to prevent from overshooting, which is important for fine-tuning large pre-trained models. For memory-efficient fine-tuning, we have applied 8-bit quantization with LoRA adapters. In this setup, we used gradient clipping (max L2 norm = 1.0) to avoid exploding gradients and keep training stable. [7]
- **Experimentation with LLaMA:** After testing the models mentioned above, we continued with LLaMA for further fine-tunings and experiments ,since it was the best-performing model. Following optimization strategies are conducted to improve our results:
 - **Layer Freezing and Unfreezing:** Initially, the transformer backbone was frozen. Later We selectively unfroze the final transformer blocks (layers 13–15) to fine-tune high-level representations without overfitting.
 - **Quantization and LoRA Fine-Tuning:** The model was loaded using 8-bit quantization via BitsAndBytesConfig, reducing memory usage while maintaining performance. To enable fine-tuning, LoRA (Low-Rank Adaptation) adapters were applied to selected layers.
 - **Threshold Optimization:** We implemented dynamic threshold tuning for each label, instead of using fixed threshold of 0.5 for converting probabilities to binary predictions. This technique searches for optimal thresholds that can maximize the F1-score for each emotion label independently. It helps to adapt to the imbalanced labels and improves final classification performance.
- **Model Evaluation Strategy:** To evaluate the model performance, we computed micro and macro F1-scores at the end of each epoch using the validation set. Additionally, we monitored the training and evaluation loss across all epochs to track learning behavior and to find potential overfitting.

VI. ANALYSIS AND INTERPRETATION

Our experiments yielded several insights:

- **BERT :** Achieved maximum macro F1 scores around 0.66 at 2nd epoch and after that there are signs of underfitting and limited sensitivity to minority classes.
- **DeBERTa:** Achieved maximum macro F1 score of 0.67 at the 9th epoch, and then the accuracy did not improve and started overfitting after that. (see Figure 2)
- **LLaMA (Full Fine-tuning):** It has a better learning curve without signs of overfitting too much and able to Achieve maximum macro F1 score of 0.67 at the 12th epoch. (see Figure 3)
- **LLaMA (Selective Unfreezing):** Our approach was to freeze all the layers of the LLaMA model except the last a few layers. We have experimented this approach with different numbers of unfrozen layers (see table II). This experiment demonstrates that, freezing all layers except the final 4 transformer layers dramatically improved the convergence and led to better results than others. In addition, Applying a 10x lower learning rate to these last layers prevents updating the parameters too much and improves the results. Despite being limited by our resources and only able to unfreeze a few layers, our approach is so effective that it achieves an F1 Macro accuracy of 0.71 by the fourth epoch. (see table I which compares the performance of different variations of LLaMA)




Fig. 3. LLaMA demonstrates a better leaning curve than DeBERTa and BERT, indicating that it has a better learning process and better fits to our problem

Model Variant	Unfrozen Layers	Macro F1	Micro F1
Llama-3.2-1B	0	0.6480	0.7117
LLaMA (3 Layers Unfrozen)	4	0.7114	0.7496
LLaMA (Quantized + LoRA)	0	0.6931	0.7421

TABLE I. PERFORMANCE COMPARISON OF LLAMA VARIANTS ON EMOTION CLASSIFICATION

VII. CONCLUSIONS AND RECOMMENDATIONS

This study demonstrates that **selective unfreezing with learning-rate separation for the unfrozen layers** is an effective technique for applying LLMs to multilabel emotion detection in low-resource scenarios. While

Model	Epoch				
	1	2	3	4	5
meta-llama/Llama-3.2-1B	0.962	0.821	0.356	0.682	0.801
with 1 layer unfreezd	0.981	0.891	0.527	0.574	0.984
with 2 layer unfreezd	0.915	0.936	0.491	0.276	0.965
with 3 layer unfreezd	0.828	0.827	0.528	0.518	0.926
with 4 layer unfreezd	0.916	0.933	0.482	0.644	0.937

TABLE II. PERFORMANCE COMPARISON OF UNFREEZING DIFFERENT LAYERS OF LLAMA MODEL

BERT-based baselines are fast and easy to train, LLaMA outperforms them when carefully fine-tuned. Our Final model with the best performance on English dataset and on Task A of the contest has the performance below:

- Macro F1: **0.7114**
- Micro F1: **0.7496**
- These results surpassed the top score on the official leaderboard (macro F1: 0.6986), although the leaderboard does not specify languages used by other participants.

Among the three evaluated models, LLaMA achieved the best performance in both macro and micro F1 scores, shows its power in generalization for frequent and rare emotion classes.

In addition, we recommend the following improvements to the contest:

Language transparency in leaderboard rankings: To fairly compare cross-language submissions it is better to have separate rankings for each language. While now there is only one leaderboard rankings that there is no information about each model about on which language it has been trained and evaluated. Some languages may have fewer number of samples so it is unfair to compare models trained on different languages. (see the comparison between their accuracy in fig 4)




Fig. 4. Comparison of macro and micro F1 scores for BERT, DeBERTa, and LLaMA with Unfreezed Layers models.

APPENDIX A

MODEL TRAINING CONFIGURATIONS AND PERFORMANCE

This appendix provides a detailed summary of the training configurations tested on each model, including batch sizes,

number of epochs, learning rates, and weight decay settings. The table also includes F1-micro and F1-macro scores for each configuration. These results highlight how different hyperparameter choices influenced performance across BERT, DeBERTa, and LLaMA models. See Table III for full details.

TABLE III. TRAINING CONFIGURATIONS AND EVALUATION RESULTS

Model	Batch Size	Epochs	LR	Weight Decay	F1 (Micro / Macro)
BERT	16	2	2e-5	—	0.70 / 0.64
BERT	16	5	2e-5	—	0.72 / 0.67
BERT	16	20	2e-5	—	0.71 / 0.66
DeBERTa	16	5	2e-5	—	0.67 / 0.64
DeBERTa	16	10	2e-5	—	0.71 / 0.65
DeBERTa	16	15	2e-5	—	0.72 / 0.67
DeBERTa	16	10	2e-5	0.01	0.70 / 0.65
LLaMA	2	10	1e-5	0.01	0.71 / 0.64
LLaMA	3	10	1e-5	0.01	0.70 / 0.63
LLaMA	3	15	1e-5	0.01	0.71 / 0.65
LLaMA	6	8	1e-5	0.01	0.74 / 0.70

REFERENCES

- [1] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding, 2019.
- [2] Pengcheng He, Xiaodong Liu, Jianfeng Gao, and Weizhu Chen. DeBERTa: Decoding-enhanced bert with disentangled attention, 2021.
- [3] Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. Lora: Low-rank adaptation of large language models, 2021.
- [4] Jiedong Lang, Zhehao Guo, and Shuyu Huang. A comprehensive study on quantization techniques for large language models, 2024.
- [5] Shamsuddeen Hassan Muhammad, Nedjma Ousidhoum, Idris Abdulmumin, Jan Philip Wahle, Terry Ruas, Meriem Beloucif, Christine de Kock, Nirmal Surange, Daniela Teodorescu, Ibrahim Said Ahmad, David Ifeoluwa Adelani, Alham Fikri Aji, Felermimo D. M. A. Ali, Ilseyar Alimova, Vladimir Araujo, Nikolay Babakov, Naomi Baes, Ana-Maria Bucur, Andiswa Bukula, Guanqun Cao, Rodrigo Tufino Cardenas, Rendi Chevi, Chiamaka Ijeoma Chukwunke, Alexandra Ciobotaru, Daryna Dementieva, Murja Sani Gadanya, Robert Geislinger, Bela Gipp, Oumaima Hourrane, Oana Ignat, Falalu Ibrahim Lawan, Rooweither Mabaya, Rahmad Mahendra, Vukosi Marivate, Andrew Piper, Alexander Panchenko, Charles Henrique Porto Ferreira, Vitaly Protasov, Samuel Rutunda, Manish Shrivastava, Aura Cristina Udrea, Lilian Diana Awuor Wanzare, Sophie Wu, Florian Valentin Wunderlich, Hanif Muhammad Zhafran, Tianhui Zhang, Yi Zhou, and Saif M. Mohammad. Brighter: Bridging the gap in human-annotated textual emotion recognition datasets for 28 languages. *arXiv preprint arXiv:2502.11926*, 2025.
- [6] Shamsuddeen Hassan Muhammad, Nedjma Ousidhoum, Idris Abdulmumin, Seid Muhie Yimam, Jan Philip Wahle, Terry Ruas, Meriem Beloucif, Christine De Kock, Tadesse Destaw Belay, Ibrahim Said Ahmad, Nirmal Surange, Daniela Teodorescu, David Ifeoluwa Adelani, Alham Fikri Aji, Felermimo Ali, Vladimir Araujo, Abinew Ali Ayele, Oana Ignat, Alexander Panchenko, Yi Zhou, and Saif M. Mohammad. SemEval-2025 task 11: Bridging the gap in text-based emotion detection. In *Proceedings of the 19th International Workshop on Semantic Evaluation (SemEval-2025)*, Vienna, Austria, 2025. Association for Computational Linguistics.
- [7] Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, Aurelien Rodriguez, Armand Joulin, Edouard Grave, and Guillaume Lample. Llama: Open and efficient foundation language models, 2023.